

```
In [1]: """ student scores analysis """
# Step-1 Importing all libraries required
import pandas as pd
import statsmodels.api as sm
from sklearn.model_selection import train_test_split
import numpy as np
import matplotlib.pyplot as plt
```

```
In [2]: # Step-2 Load dataset into dataframe
student_score_df = pd.read_csv("students_scores.csv")
print('First 10 entries: \n', student_score_df.head(10))
```

First 10 entries:

	Hours	Scores
0	2.5	21
1	5.1	47
2	3.2	27
3	8.5	75
4	3.5	30
5	1.5	20
6	9.2	88
7	5.5	60
8	8.3	81
9	2.7	25

```
In [3]: print('Finding missing value: \n', student_score_df.isnull().sum())
```

Finding missing value:

```
Hours      0
Scores     0
dtype: int64
```

```
In [4]: print('Information of database: \n', student_score_df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 25 entries, 0 to 24
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Hours   25 non-null        float64
1   Scores  25 non-null        int64
dtypes: float64(1), int64(1)
memory usage: 528.0 bytes
Information of database:
None
```

```
In [5]: student_score_df.describe()
```

```
Out[5]:
```

	Hours	Scores
count	25.000000	25.000000
mean	5.012000	51.480000
std	2.525094	25.286887
min	1.100000	17.000000
25%	2.700000	30.000000
50%	4.800000	47.000000
75%	7.400000	75.000000
max	9.200000	95.000000

```
In [6]: # Step-3 Creating feature variable(X) and outcome variable(y) for building model
data = student_score_df[["Hours", "Scores"]]
predict = "Scores"
X1 = np.array(data.drop([predict], 1))
X = sm.add_constant(X1)
Y = np.array(data[predict])
print('Values of X: \n', X)
print('Values of Y: \n', Y)
```

Values of X:

```
[[1.  2.5]
 [1.  5.1]
 [1.  3.2]
 [1.  8.5]
 [1.  3.5]
 [1.  1.5]
 [1.  9.2]
 [1.  5.5]
 [1.  8.3]
 [1.  2.7]
 [1.  7.7]
 [1.  5.9]
 [1.  4.5]
 [1.  3.3]
 [1.  1.1]
 [1.  8.9]
 [1.  2.5]
 [1.  1.9]
 [1.  6.1]
 [1.  7.4]
 [1.  2.7]
 [1.  4.8]
 [1.  3.8]
 [1.  6.9]
 [1.  7.8]]
```

Values of Y:


```
[21 47 27 75 30 20 88 60 81 25 85 62 41 42 17 95 30 24 67 69 30 54 35 76
 86]
```

```
In [7]: # Step-4 Splitting dataset into Training and Validation Sets
train_X, test_X, train_y, test_y = train_test_split(X, Y, train_size=0.8, random_state=0)
```

```
In [8]: # Step-5 Fitting of regression model by OLS method
stu_scores_lm = sm.OLS(train_y, train_X).fit()
print('Parameters: ', stu_scores_lm.params)
line = stu_scores_lm.params[1]*X1+stu_scores_lm.params[0]
```

Parameters: [2.01816004 9.91065648]

```
In [9]: # Plotting for the test data
plt.scatter(X1, Y)
plt.plot(X1, line)
plt.title('Hours Vs Scores', fontsize=14)
plt.xlabel('Hours', fontsize=14)
plt.ylabel('Scores', fontsize=14)
plt.grid(True)
plt.show()
```



```
In [10]: # Statistical data required for diagnosing regression model
print('Summary of model:', stu_scores_lm.summary2())
```

```
Summary of model:                      Results: Ordinary least squares
=====
```

Model:	OLS	Adj. R-squared:	0.949
Dependent Variable:	y	AIC:	129.3715
Date:	2021-07-15 01:47	BIC:	131.3630
No. Observations:	20	Log-Likelihood:	-62.686
Df Model:	1	F-statistic:	353.5
Df Residuals:	18	Prob (F-statistic):	2.79e-13
R-squared:	0.952	Scale:	34.331

```
-----
```

	Coef.	Std.Err.	t	P> t	[0.025	0.975]
const	2.0182	3.0570	0.6602	0.5175	-4.4043	8.4407
x1	9.9107	0.5271	18.8023	0.0000	8.8033	11.0181

```
-----
```

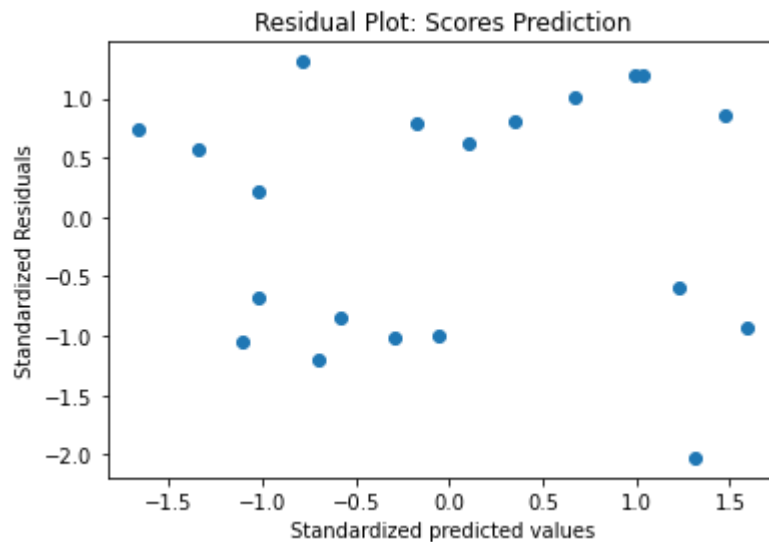
Omnibus:	4.659	Durbin-Watson:	1.813
Prob(Omnibus):	0.097	Jarque-Bera (JB):	1.720
Skew:	-0.296	Prob(JB):	0.423
Kurtosis:	1.691	Condition No.:	14

```
=====
```

```
In [11]: # Step-6 Model Dignostics to validate the data
# Residual analysis-Test of Homoscedasticity
stu_scores_resid = stu_scores_lm.resid
```

```
In [12]: def get_standardized_values(vals):
        """
        Test of Homoscedasticity
        :param vals: series of variable values
        :return: standardized values of vals
        """
        return (vals - vals.mean())/vals.std()
```

```
In [13]: plt.scatter(get_standardized_values(stu_scores_lm.fittedvalues), get_standardi
zed_values(stu_scores_resid))
plt.title("Residual Plot: Scores Prediction")
plt.xlabel("Standardized predicted values")
plt.ylabel("Standardized Residuals")
plt.show()
```



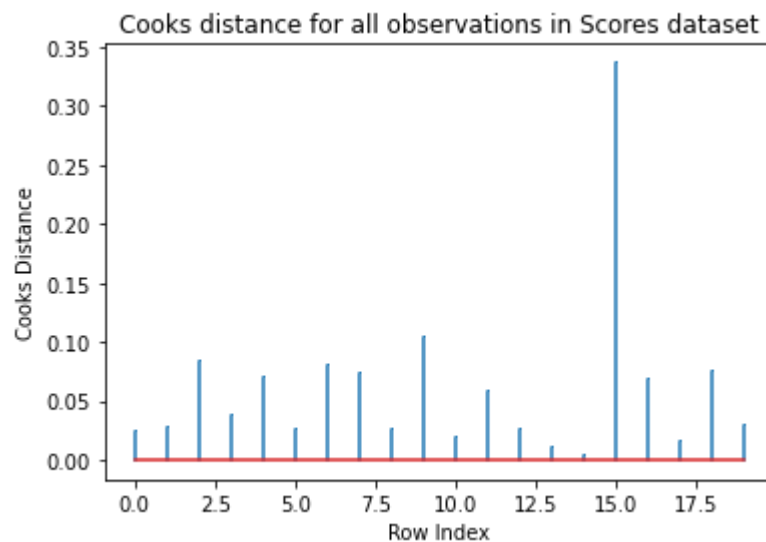
```
In [14]: # Outlier Analysis(most influential observation)- Z_score
from scipy.stats import zscore
student_score_df['z_score_percent'] = zscore(student_score_df.Scores)
Z_score = student_score_df[(student_score_df.z_score_percent > 3.0) | (student
_score_df.z_score_percent < -3.0)]
print(Z_score)
```

```
Empty DataFrame
Columns: [Hours, Scores, z_score_percent]
Index: []
```

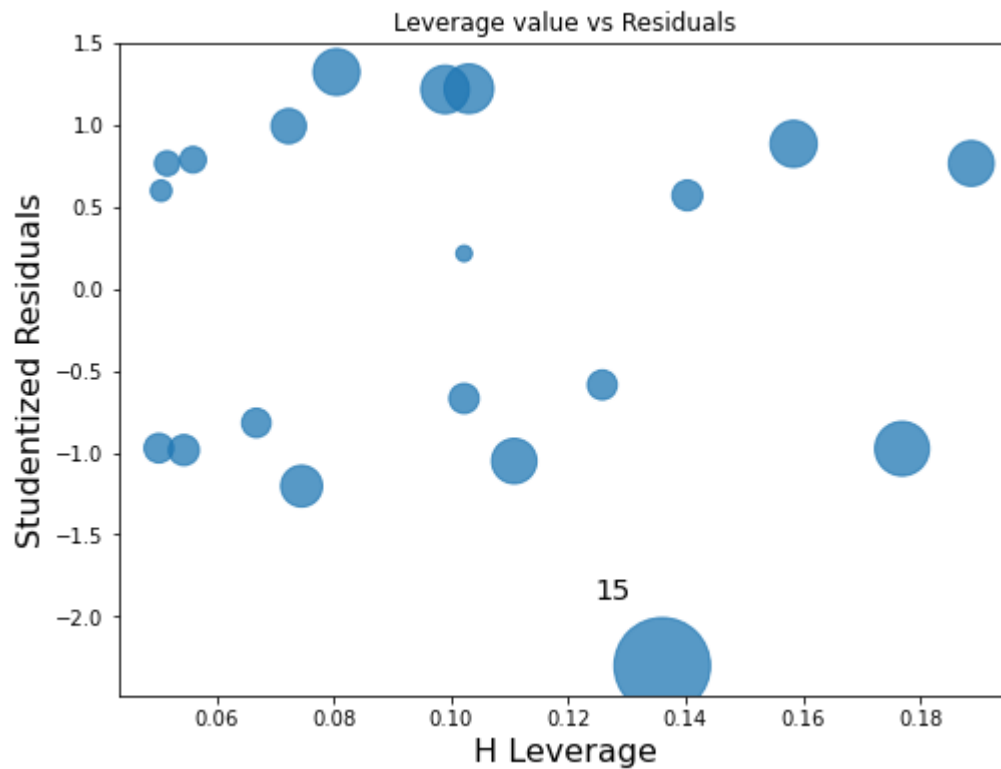
```
In [15]: # Cook's Distance
score_influence = stu_scores_lm.get_influence()
(c, p) = score_influence.cooks_distance
plt.stem(np.arange(len(train_X)),
         np.round(c, 3),
         markerfmt=",")
plt.title("Cooks distance for all observations in Scores dataset")
plt.xlabel("Row Index")
plt.ylabel("Cooks Distance")
plt.show()
```

<ipython-input-15-6238179c4d94>:4: UserWarning: In Matplotlib 3.3 individual lines on a stem plot will be added as a LineCollection instead of individual lines. This significantly improves the performance of a stem plot. To remove this warning and switch to the new behaviour, set the "use_line_collection" keyword argument to True.

```
plt.stem(np.arange(len(train_X)),
```



```
In [16]: # Leverage values
from statsmodels.graphics.regressionplots import influence_plot
fig, ax = plt.subplots(figsize=(8, 6))
influence_plot(stu_scores_lm, ax=ax)
plt.title("Leverage value vs Residuals")
plt.show()
```



```
In [17]: # Step-7 Making Prediction
# predicting using validation set
pred_y = stu_scores_lm.predict(test_X)
# print(pred_y)
```

```
In [18]: df = pd.DataFrame({'Actual': test_y, 'Predicted': pred_y})
print(df)
```

	Actual	Predicted
0	20	16.884145
1	27	33.732261
2	69	75.357018
3	30	26.794801
4	62	60.491033

```
In [19]: # you have to create a DataFrame since the Statsmodels formula interface expects it

new_X = pd.DataFrame({'Const': [1], 'Hours': [9.25]})

y_new = stu_scores_lm.predict(new_X)

# Y = np.array(data[predict])

print("Given out of sample input hours : \n", new_X)
print("Prediction of scores : \n", y_new)
```

```
Given out of sample input hours :
   Const  Hours
0      1   9.25
Prediction of scores :
0  93.691732
dtype: float64
```

```
In [20]: # Formula for Mean Absolute Percentage Error
def mean_absolute_percentage_error(y_true, y_pred):
    y_true, y_pred = np.array(y_true), np.array(y_pred)
    return np.mean(np.abs((y_true - y_pred) / y_true)) * 100
```

```
In [21]: # Step-8 Different measures for accuracy of prediction

from sklearn import metrics
print('Mean Absolute Error: \t \t \t',
      metrics.mean_absolute_error(test_y, pred_y))

from sklearn.metrics import r2_score, mean_squared_error
print('r2_score: \t \t \t \t', np.abs(r2_score(test_y, pred_y)))

print('Mean Squared Error: \t \t \t', np.sqrt(mean_squared_error(test_y, pred_y)))

print('Mean Absolute Percentage Error: \t', mean_absolute_percentage_error(test_y, pred_y))
```

```
Mean Absolute Error:          4.183859899002978
r2_score:                    0.9454906892105355
Mean Squared Error:          4.647447612100368
Mean Absolute Percentage Error: 12.568891617045674
```